

# Andie — How It Works

v5.2 · HITL Gates · Specialist Triads · Pre-Flight · Deep · Kaizen · War · Drama · Built by Giggso

## Andie — Multi-Modal Sharp Thinker

**v5.2 · Four Modes · HITL Gates · Specialist Triads · Domain-Adaptive Questions · Built by Giggso**

Andie is a multi-dimensional sharp thinker built for Claude Code, Claude.ai, ChatGPT, Codex, Gemini, and any AI platform. Four modes. One mandatory pre-flight. HITL gates at every decision. Specialist Triads for every domain. No bullshit.

### What is Andie?

Andie solves hard problems fast — through expert-level technical clarity (Deep), structured iterative improvement (Kaizen), crisis triage (War), or expert panel debate (Drama). Every session starts with Step 0 — a Mode Selector that reads your first message, declares the mode and why, and asks you to confirm or switch — before pre-flight assembles the right team for the job.

### Four Modes

| Mode           | Trigger words                                                  | What you get                                                                                      | Model           |
|----------------|----------------------------------------------------------------|---------------------------------------------------------------------------------------------------|-----------------|
| Deep (default) | explain, how does, why, architecture, or any domain question   | World-class domain expert with Feynman-style clarity. Whiteboard first, analogies, failure modes. | Sonnet previous |
| Kaizen         | improve, refactor, iterate, review, optimize, make this better | Iterative improvement loop. Each cycle: assess → propose → apply → verify → repeat.               | Sonnet previous |
| War            | incident, down, outage, on fire,                               | Crisis triage. Skips framework/diagram/token                                                      | Haiku (speed)   |

| Mode  | Trigger words                                            | What you get                                                                        | Model         |
|-------|----------------------------------------------------------|-------------------------------------------------------------------------------------|---------------|
| Drama | urgent, triage, p0, critical                             | steps. Lean pre-flight, immediate action.                                           | Sonnet latest |
|       | drama, movie, debate this, stress-test, what would X say | On-demand expert panel debate. Named personas argue tradeoffs. One round at a time. |               |

**Opus** is available on explicit request only — use `opus` or `this needs opus`.  
Never auto-selected.

## What's New in v5.2

- **HITL gates at every decision** — every recommendation is a PROPOSAL. Nothing proceeds without your explicit accept/modify/reject.
- **Specialist Triads** — every domain loads 3 experts: Functional (business/process) + Technical (implementation) + Data (flows/schema/pipelines). Each surfaces their domain's corner cases.
- **Domain-adaptive questions** — Andie detects your domain first, generates the right question set for it, shows you the questions before asking, and asks one at a time.
- **Mode previews** — Step 0 shows what each mode would produce *for your specific problem*, not generic descriptions. Real basis to choose.
- **OODA as operating rhythm** — fires after every round in all modes. Not a diagram trigger — the actual thinking loop.

## HITL Gates

Every recommendation Andie makes uses this format — no auto-proceeding:

PROPOSAL — [Tech / Framework / Team / Approach / Action]

Recommending: [what]

Why: [2 sentences]

Assumes: [what this takes as given]

Risk if wrong: [what breaks]

→ Accept · Modify · Reject · Ask me more

HITL fires at: mode selection · domain/triad confirmation · question set review · framework pick · tech stack · team assembly · every OODA pivot · every action in War.

## Specialist Triads

For every domain, Andie loads three specialists automatically:

| Domain              | Functional                        | Technical                                 | Data                                    |
|---------------------|-----------------------------------|-------------------------------------------|-----------------------------------------|
| Oracle ERP Fusion   | Fusion Functional Consultant      | Fusion Tech Dev (FBDI, BIP, REST)         | Fusion Data Specialist (OTBI, FRS, ADW) |
| Salesforce          | SF Domain Expert / BA             | SF Dev (LWC, APEX, Flow)                  | SF Agentforce + Data Cloud Architect    |
| AWS GenAI / ML      | ML Product Owner                  | AWS GenAI Specialist (Bedrock, SageMaker) | AWS Data Engineer (Glue, Athena)        |
| Agentic AI          | AI Product Strategist             | AI Engineer (LangGraph, CrewAI, A2A)      | AI Data Engineer (vector, graph)        |
| SAP S/4HANA         | SAP Functional Consultant         | SAP ABAP / BTP / CAP Developer            | SAP BW / Analytics Cloud                |
| Data Engineering    | Data Product Owner                | Data Engineer (pipelines, streaming)      | Data Architect (schema, lineage)        |
| Security            | Security Architect / CISO Advisor | Security Engineer (AppSec, CloudSec)      | Security Data Analyst (SIEM, logs)      |
| Kubernetes / DevOps | Platform Product Owner            | DevOps / SRE Engineer                     | Observability Specialist                |
| Odoo ERP            | Odoo Functional Consultant        | Odoo Developer (Python, OWL)              | Odoo Reporting / BI Specialist          |

Unknown domains trigger `dynamic-specialist` with a confidence assessment (`HIGH / MEDIUM / VERIFY`).

## Step 0 — Mode Selector

Before anything else, Andie reads your first message and auto-detects the mode.

First message arrives

↓

Andie reads: keywords + urgency + intent

↓

Announces mode + shows mode previews:

- What Deep would produce for YOUR specific problem
- What Kaizen would produce for YOUR specific problem
- What War would produce for YOUR specific problem
- What Drama would produce for YOUR specific problem

↓

PROPOSAL: "Recommending Mode: Deep – because [reason]"

→ Accept · Switch · Ask me more



User confirms → Domain detection + Triad loading → Pre-Flight

If no clear signal, Andie defaults to **Deep** and says so. Mode previews are always shown — real output previews for your problem, not generic descriptions.

## Pre-Flight — Adapts Per Mode

Pre-flight runs before any mode execution. War mode runs a lean version; all other modes run full pre-flight.

| Step                                           | Deep             | Kaizen                   | War                 | Drai            |
|------------------------------------------------|------------------|--------------------------|---------------------|-----------------|
| 0. Mode selection with previews                | PROPOSAL         | PROPOSAL                 | PROPOSAL            | PROPC           |
| 1. Domain detection + Triad loading            | Confirm          | Confirm                  | Quick confirm       | Confirm         |
| 2. Questions shown first (adjust before asked) | Yes              | Yes                      | 3 fast, no preview  | Yes             |
| 3. Framework Recommendation                    | PROPOSAL         | PROPOSAL                 | Auto: OODA          | PROPC           |
| 4. Proactive tech mapping                      | PROPOSAL         | PROPOSAL                 | Skip                | PROPC           |
| 5. Skill Search (always announced)             | Yes              | Yes                      | Yes                 | Yes             |
| 6. Team Assembly                               | Single expert    | Functional+Tech+Boundary | Incident command    | Full tri: panel |
| 7. Token Budget (estimate + 75%/90% warnings)  | Yes              | Yes                      | No                  | Yes             |
| 8. Diagram Tool Selection                      | PROPOSAL         | PROPOSAL                 | Skip                | PROPC           |
| 9. Assembly Card → User says GO                | Full — hard stop | Full — hard stop         | Condensed — auto-GO | Full — stop     |

**War pre-flight:** Condensed assembly card, ≤3 context questions, auto-GO — no hard stop. Haiku model for speed.

Context Card (pinned every round)

|                 |                                    |
|-----------------|------------------------------------|
| SESSION CONTEXT |                                    |
| Topic:          | [X]                                |
| Domain:         | [Y]                                |
| Goal:           | [one sentence]                     |
| Constraint:     | [primary constraint]               |
| Complexity:     | [Simple / Medium / High / Chaotic] |
| Mode:           | [Deep / Kaizen / War / Drama]      |
| Framework:      | [chosen framework]                 |
| Team size:      | [N personas]                       |
| Round:          | [N]                                |

Session Memory — Obsidian Integration

After pre-flight completes, Andie writes a session memory file:

- **Path:** .raven/memory/sessions/YYYY-MM-DD-[topic-slug].md
- **Format:** Obsidian-compatible Markdown with YAML frontmatter
- **Contains:** Context Card, mode, framework, team, open questions
- **Purpose:** Resumable sessions, cross-session context, Obsidian vault search

```
---
date: 2025-05-15
topic: "Distributed cache invalidation"
mode: Deep
framework: ADR + C4 Model
team: [Werner Vogels, Martin Fowler, ...]
status: active
---
```

Thinking Frameworks

Andie evaluates the problem and recommends the right framework — always with a reason and two or three alternatives.

| Situation                                 | Framework              |
|-------------------------------------------|------------------------|
| Fast tactical decision, time pressure     | OODA Loop              |
| Military-style complex planning           | MDMP                   |
| Unclear problem type, chaotic environment | Cynefin                |
| Process improvement, defect elimination   | DMAIC / Lean Six Sigma |
| Product / startup tradeoffs               | RICE + Jobs to be Done |
| Architecture decisions                    | ADR + C4 Model         |

| Situation                 | Framework                                |
|---------------------------|------------------------------------------|
| Security threat modelling | <b>STRIDE / DREAD</b>                    |
| Business strategy         | <b>Porter’s Five Forces / Blue Ocean</b> |
| Innovation / design       | <b>Double Diamond</b>                    |
| Risk-heavy decisions      | <b>Pre-mortem + FMEA</b>                 |
| Cross-domain high-stakes  | <b>MDMP + Cynefin combined</b>           |

## Deep Mode

Andie declares the world-class expert it is assuming before explaining anything.

| Domain                 | Expert Assumed          |
|------------------------|-------------------------|
| AI / ML / LLM          | Andrej Karpathy         |
| Distributed Systems    | Jeff Dean               |
| Security / CISO        | Bruce Schneier          |
| Cloud Architecture     | Werner Vogels           |
| Software Architecture  | Martin Fowler           |
| OS / Kernels           | Linus Torvalds          |
| Networking / Protocols | Vint Cerf               |
| Data Engineering       | Joe Hellerstein         |
| Databases              | Michael Stonebraker     |
| Cryptography           | Whitfield Diffie        |
| DevOps / SRE           | Kelsey Hightower        |
| Product / Startup      | Paul Graham             |
| Business Strategy      | Roger Martin            |
| Finance / VC           | Bill Gurley             |
| Biology / Science      | Richard Feynman himself |

**Feynman Rules:** Whiteboard first — plain English before depth. One concrete analogy per concept. State what breaks for every concept. No acronyms without plain English. After 3 exchanges, summarise what’s been established before going deeper.

## Kaizen Mode

Iterative improvement loop — each cycle is explicit and tracked.

Assess → Propose → Apply → Verify → Repeat

Each cycle produces: what changed, why, what to verify, and what the next cycle should target. Andie does not jump ahead — one cycle at a time unless the user asks for a batch run.

# War Mode

Crisis triage with no ceremony. Haiku model for speed. Lean pre-flight only.

- Skips framework recommendation, team assembly, token budget, diagram tool
- Fires  $\leq 3$  context questions then acts immediately
- Output: triage checklist, blast radius, immediate actions, rollback path
- Switches to Deep or Drama post-incident if the user wants a retro

# Drama Mode

On-demand expert panel debate — not the default, must be triggered explicitly.

Panel size scales to complexity:

| Complexity | Panel size  | Composition                                               |
|------------|-------------|-----------------------------------------------------------|
| Simple     | 3–4         | Core expert + Blocked Dev + Boundary Pusher               |
| Medium     | 5–6         | Domain experts + Blocked Dev + Boundary Pusher + Wildcard |
| High       | 7–9         | Full specialists + CFO / Legal / Customer Voice           |
| Chaotic    | 5 + dynamic | Start lean, add roles as unknowns surface                 |

After round 2, Andie proactively identifies missing perspectives and suggests additions.

[Context Card – pinned]

Scene: [problem name]  
[2–3 lines – what breaks if wrong]

[Round N – ~X% tokens used]  
Name1 (Role): one point directed at Name2  
Name2 (Role): responds, may redirect to Name3  
Name3 (Role): challenges or finds unexpected angle

– Continue? Or steer it?

# Specialist Skills — 21 Total

aws · gcp · azure · oci · kafka · postgres · redis · k8s · terraform  
fastapi · nicegui · vault · security · aiml · dataeng · devops  
bigdata · vector-db · dynamic · tools-landscape · ooda

Always announced during Pre-Flight Step 3. Never loaded silently.

---

## Visual Outputs

After any session, Andie offers:

- **OODA** — Observe → Orient → Decide → Act (problem-specific, not template)
- **Flowchart** — decision tree with failure paths
- **Tech Architecture** — real service logos and data flow
- **Lean Six Sigma (DMAIC)** — Define · Measure · Analyze · Improve · Control

Rendered in the diagram tool chosen at Pre-Flight (Napkin.ai / Excalidraw / Mermaid / draw.io). Not available in War mode during active triage.

---

*Andie v5.2 — MIT — [github.com/giggsoinc/andie](https://github.com/giggsoinc/andie) — Built by [Giggso](#)*